


Who answers for *an AI agent?*



TRUSTLAYER
FOUNDATION




An identity protocol for AI agents, and the nonprofit that stewards it.

THE PROBLEM


Agents already act.

When something goes wrong, there is no standard way to establish who answers for it.



This is not an authentication problem. It is an accountability problem.

We are not the first to say this.



**The UK DRCF foresight report, Singapore's IMDA, the Australian Productivity Commission, the WE BUILD non-paper (Bosch, Google Payments, Visa Europe, Ericsson, Spherity), Birch and Hoffart on Know Your Agent, and GLEIF's own August working paper*

One layer is missing.

Agent interoperability and tool connectivity

Agent-to-service authentication

Who answers for it.

NO STANDARD

Identity inside a single trust domain

Transport, DNS, PKI

The identified person or organization behind the agent, verifiable by a third party with no prior relationship to anyone involved.

An identity that only works inside one provider is not an identity. It is an account.



Who answers for this agent?

ARIA gives every AI agent a verifiable identity. End to end, in four steps.



Registered

STEP 01 · REGISTRY

A registry verifies the human and the organization behind the agent — as far as the level demands — and issues its credential: the AID.



Signed & anchored

STEP 02 · CREDENTIAL

The AID is a W3C credential with a post-quantum signature. From L1, it is anchored to the holder's domain in DNS — the trust infrastructure that already exists.



Verified locally

STEP 03 · VERIFICATION

Signature, validity and level are checked in-process; the level is derived from evidence, never declared. Revocation: a cached Status List, refreshed within 60 s. No registry call when deciding.



Admitted, or not

STEP 04 · ATP

The receiver publishes what it requires in DNS and admits or refuses locally. DMARC for email; ATP for agents.


The Foundation stewards the standard and signs. Registries verify. **Nobody sees personal data that isn't theirs to see.**

***Disambiguation note for W3C CG. Different from WAI-ARIA**



Four questions. Four levels.

Each level adds a branch of evidence over the same cryptographic root.



L2 proves the agent acts WITH the organization. · L3 proves someone can BIND it.



The identity arrives at the door. *Now what?*

ARIA gives the agent an identity. ATP is what the receiver does when that identity shows up.

01

DECLARE

Publish your admission conditions in a DNS record. Public, machine-readable and optional.

```
_aria-policy    TXT    "v=ATP1; min=L2; req=finance:*
```

02

EVALUATE

Check on your own machine: signature, validity, revocation status, level and scopes.

03

ADMIT


Or deny with a code that says exactly what was missing. No emails, no tickets.

ATP-DENY-LEVEL · ATP-DENY-SCOPE

Nobody registers with anybody. The receiver decides locally, with no arbiter in between, and an unknown agent can tell whether it qualifies before calling.

¿REMEMBER THE PROBLEM?

The identity arrives at the door. *Now what?*



Three things worth your scrutiny.

01

The level is derived, not declared

Two independent axes: person and entity. Any verifier recomputes the level from the evidence. A label that does not match its facts is a malformed credential.

02

Local evaluation, zero coordination

Admission goes through no one: neither the foundation nor the issuing registry sits in the hot path. It works between strangers, which is precisely the hard case.

03

Post-quantum from birth

ML-DSA-65 (FIPS 204) and Ed25519 (RFC 8032), composite, with AND logic: if either half fails, the credential is invalid. The launch configuration, not a migration plan.

NOT A WHITE PAPER



Public specification. Open code.

Everything described here is published and verifiable today.

THE AGENT IDENTIFIER

```
did:aria:aria.bar:agent-alpha
```

VERIFICATION, PUBLIC AND KEYLESS

```
curl https://api.aria.bar/v1/verify/did:aria:...
```

THE RECEIVER POLICY, IN DNS

```
_aria-policy TXT "v=ATP1; min=L2; req=finance:*
```

Public repository

github.com/trustlayer-foundation

Specification

aria.bar/spec

Licensing

CC BY 4.0 · Apache 2.0

Anyone can verify a credential without asking us, and without going through us.



Built on current work.

ARIA — cross-organizational identity layer

W3C Decentralized Identifiers

the agent identifier

W3C Verifiable Credentials 2.0

the credential format

W3C Bitstring Status List

revocation propagation

DNS · FIPS 204 · RFC 8032

anchoring and signatures

*We are already downstream
of official Recommendations.*

We come with implementation findings, not a parallel universe.

ARIA complements agent interoperability protocols, workload identity, cloud IAM and the emerging IETF frameworks. It replaces none of them.



What ARIA does *not* solve.

- **Behavior**

It asserts verified facts at a date, never future conduct.
- **Internal identity**

ARIA sits at the door, not inside your network.
- **Full legal capacity**

Age of majority only. Guardianship is not federably queryable.
- **Injection attacks**

The presentation-attack standard does not cover them. Separate standard, already flagged.
- **Delegation**

Agent-to-agent delegation is not specified yet.

We would rather show you the holes we have found than the demo we have polished.

The trust layer *nobody owns.*

A nonprofit. Not part of the operation: there is nothing for it to hold.



Holds no private keys

neither the agents' nor the registries'



Sees no personal data

no documents, no biometrics, no case files



Does not operate the registry

verification is done by independent entities

THE ONE THING IT DOES

Steward the standard. And the standard is open.

CC BY 4.0 on the specification. Apache 2.0 on the code. A public RFC process. Anyone can implement it without asking permission.

If trust can be owned, it is not trust.

Anti-capture by construction.

GOVERNANCE DNA

Dual-entity model

nonprofit and commercial as siblings

Meritocracy, lazy consensus

community-driven development

Working groups, IPR policy

structured contribution

Rough consensus, running code

proven in this room

LICENSING AND IPR

CC BY 4.0

all specifications

Apache 2.0

all code

DCO over CLA

on every contribution

Royalty-free patents

perpetual, irrevocable license

No spec change may advantage a specific implementation. No standard may require proprietary infrastructure.

WHO WE ARE



The people behind.



Aaron Grego

CEO

Direction of the foundation and the standard.



Iván Moreno

CTO

Protocol architecture and implementation.




Adolfo Grego

**SECRETARY ·
INSTITUTIONAL RELATIONS**

Governance, legal framework and outreach.

Punto 2012 - Registry Operators over 14 years

TrustLayer Foundation A.C. · Constituted March 2026 · Mexico City



TRUSTLAYER
FOUNDATION



Thank you!

What now?

Where do we go from here?